

Hyperparameters

Settings, values, and their justification

Why every hyperparameter has the value it does, how it was chosen, and where the value comes from. Values are the production settings in `configs/*.yaml` and the classifier code. The formulas these feed are independently verified against the code by `scripts/verify_math.py` (all checks pass — see the end of this doc).

Three sources of a value are used throughout:

- **Architecture-fixed** — set by the supervisor-mandated design (IWSCN reference paper); not free to change.
- **Theory/standard** — a standard recipe from the literature (cited).
- **Empirically chosen** — selected on the validation split / by cross-validation (never on test).

1. Wavelet-Scattering branch (`configs/model.yaml` → `model.scattering`)

Param	Value	Why this value	How chosen
J (scale)	2	Each order-1 coefficient downsamples by 2^J . $J=2 \rightarrow 64 \times 64$ feature maps from 256×256 , keeping enough spatial resolution to localise small defects; $J=3$ would drop to 32×32 and blur tiny porosity.	Architecture-fixed (IWSCN); confirmed adequate in the type-classifier search (<code>scripts/tune_type_classifier.py</code> tried $J=3$, no gain).
L (orientations)	8	8 wavelet orientations (every 22.5°) capture directional weld texture (slag is elongated, porosity round).	Theory (standard Kymatio default; Bruna & Mallat).
order	2	Order-2 scattering captures interactions between scales — needed to separate texture classes; order-1 alone is too weak.	Architecture-fixed (IWSCN).

Coefficient count (verified): $1 + J \cdot L + L^2 \cdot J(J-1)/2 = 1 + 16 + 64 = 81$. This is the number of scattering channels, checked against the code in `scripts/model_spec.py`.

2. Network architecture (`configs/model.yaml` → `model`)

Param	Value	Why this value	How chosen
base	32	Encoder width 32→64→128→256, bottleneck 512 → 8.8 M params. Large enough to learn weld texture, small enough to train on a CPU/free GPU with 1050 images without heavy overfitting.	Empirically (standard U-Net width that fits memory).
seg_classes	3	background / porosity / slag — per-pixel segmentation needs a background class.	Architecture-fixed (decision §3).
n_classes	2	porosity vs slag image-level type.	Architecture-fixed (decision §2).
dropout	0.1	Light regularisation; higher values hurt the already-small defect signal.	Empirically.
num_groups	8	GroupNorm groups. GroupNorm (not BatchNorm) because batch size is tiny (4) — BatchNorm statistics would be noisy. 8 groups is the standard GN default.	Theory (Wu & He, GroupNorm).
size_head	false	Defect size is measured geometrically from the mask (regionprops), which is exact, instead of regressed (which would need size labels we don't have).	Architecture-fixed (decision §11).

3. Loss weights (configs/model.yaml → loss)

$L = w_{\text{seg}} \cdot L_{\text{seg}} + w_{\text{cls}} \cdot \text{CE} + w_{\text{boundary}} \cdot \text{boundary}$, with $L_{\text{seg}} = \text{dice} \cdot \text{Dice} + \text{focal} \cdot \text{FocalCE} + \text{tversky} \cdot \text{FocalTversky}$.

Param	Value	Why this value
w_seg	0.6	Segmentation (detection) is the primary task → largest weight.
w_cls	0.3	Classification is secondary (type is finally decided by the scattering classifier).
w_boundary	0.1	A small nudge for crisp defect edges; too high distracts from region overlap.

Param	Value	Why this value
tversky_alpha	0.3	Weight on false positives .
tversky_beta	0.7	Weight on false negatives . $\beta > \alpha$ ($0.7 > 0.3$) makes <i>*missed*</i> defect pixels hurt more than false alarms — correct for thin/faint defects that fail by being missed.
tversky_gamma	1.0	Focal-Tversky exponent. 1.0 = plain Tversky; left at 1.0 because $\beta > \alpha$ already handles the imbalance and higher γ destabilised training.
focal_gamma	2.0	Focal-CE focusing parameter; 2.0 is the standard value (Lin et al., Focal Loss) — down-weights easy background pixels.
include_background_in_dice	false	Critical: averaging Dice over foreground only. With background included the model scored ~ 0.66 Dice by predicting <i>*all background*</i> and abandoning the defect (the collapse bug). Excluding background removes that free credit.

Why these weights and not others: the split 0.6 / 0.3 / 0.1 follows the task priority (detect > classify > refine edges); the Tversky α/β follow the standard imbalance recipe (Salehi et al.); the include_background=false choice was forced by the observed collapse.

4. Training schedule (configs/model.yaml → train)

Param	Value	Why this value	How chosen
epochs	150	Enough for the cosine schedule to converge; early-stopping ends runs sooner if val Dice plateaus.	Empirically.
batch_size	4	Largest batch that fits in CPU/free-GPU memory at 256×256; small batch is <i>*why*</i> we use GroupNorm.	Hardware-constrained.
lr	0.001	Standard Adam starting LR; with warmup + cosine it anneals to ~ 0 . Higher diverged, lower was slow.	Empirically (standard default).
weight_decay	1e-4	Light L2 regularisation, standard for Adam-family on small data.	Theory/standard.

Param	Value	Why this value	How chosen
warmup_epochs	5	Avoids large early steps destabilising the fresh weights before cosine decay.	Standard recipe.
scheduler	cosine	Smooth anneal → better final minima than step decay on small data.	Standard recipe.
grad_clip	1.0	Caps gradient norm so the focal/Tversky terms can't spike the update.	Standard recipe.
early_stop_patience	25	Stop if val foreground-Dice hasn't improved in 25 epochs — guards against overfitting the tiny set.	Empirically.
class_weighting	true	Inverse-frequency weighting on the classification CE. The dataset is now 525/525 balanced, so it is effectively neutral — kept as a safeguard.	Safeguard.

Augmentation (configs/model.yaml → augment, train split only): h/v flips & 90° rotations (p=0.5) are valid because defect appearance is orientation-agnostic; gain (0.8–1.2), speckle ($\sigma=0.05$) and attenuation (0.7) simulate real PAUT acquisition variation so the model is robust to scanner/gain differences.

5. Type classifier (src/models/type_classifier.py)

Param	Value	Why this value	How chosen
pipeline	log1p → StandardScaler → MLP(64)	log1p is the standard *log-scattering* renormalisation (Mallat/Bruna); StandardScaler centres features; a tiny 64-unit MLP is the "small classifier" of the IWSCN recipe.	5-fold CV on train (tune_type_classifier.py).
MLP hidden	64	Big enough to separate the two texture classes, small enough not to overfit the small training set.	CV-selected over LogReg / SVM / MLP.
alpha (L2)	1e-2	MLP weight decay; CV-tuned for best balanced accuracy.	CV.
result	0.88 balanced (vs 0.80 plain LogReg)	—	CV 0.87 ≈ test 0.88 → generalises.

6. Severity & digital-twin (configs/characterize.yaml, configs/twin.yaml)

*All values here are **tunable placeholders, NOT validated acceptance criteria** — they must be calibrated to ISO 5817 / ASME BPVC with a domain expert before any real pass/fail use.*

Param	Value	Status
severity.type_weight porosity/slag	1.0 / 1.2	Placeholder (slag weighted higher as generally more critical).
severity.size_ref_px	200	Placeholder area→severity reference.
twin.capacity	5.0	Placeholder severity "budget" in health = $1 - \min(1, \Sigma \text{severity} / \text{capacity})$.
twin.review_health	0.7	Placeholder: health below this → REVIEW.
cutoffs moderate / critical	0.5 / 1.0	Placeholder severity bands.

7. Math == code verification

scripts/verify_math.py recomputes each formula in independent numpy and asserts it matches the production code (tolerance 1e-6). On a hand-computable 2×2 / 3-class example:

Quantity	By hand / numpy	Code	Match
Foreground Dice loss	0.407895	0.407895	✓
Focal-Tversky loss ($\alpha 0.3, \beta 0.7$)	0.414948	0.414948	✓
Focal-CE loss ($\gamma 2$)	0.079805	0.079805	✓
Health index ($\Sigma \text{sev } 1.2, \text{ cap } 5$)	0.760	0.760	✓
Health index ($\Sigma \text{sev } 4.5, \text{ cap } 5$)	0.100	0.100	✓

Reproduce: `py -3.14 -m scripts.verify_math` (exits 0 = all match).